

DS6030 Project Data Evaluation

Tyler Hobbs

2025-06-17

```
library(tidyverse)
library(GGally)
library(tidymodels)
library(discrim)
library(patchwork)
library(probably)
library(vip)
```

Loading the data and Pre-processing

```
train<-read_csv("HaitiPixels.csv")
```

```
## Rows: 63241 Columns: 4
## — Column specification —————
## Delimiter: ","
## chr (1): Class
## dbl (3): Red, Green, Blue
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```

test1<-read_table("orthovnir057_R0I_NON_Blue_Tarps.txt",show_col_types = FALSE)
test1_sub<-test1[1:10]
colnames(test1_sub)<-c("ID","X","Y","Map X","Map Y", "Lat","Lon","B1","B2","B3")
test2<-read_table("orthovnir067_R0I_Blue_Tarps_data.txt",show_col_types = FALSE)
test3<-read_table("orthovnir067_R0I_Blue_Tarps.txt",show_col_types = FALSE)
test3_sub<-test3[1:10]
colnames(test3_sub)<-c("ID","X","Y","Map X","Map Y", "Lat","Lon","B1","B2","B3")
test4<-read_table("orthovnir067_R0I_NOT_Blue_Tarps.txt",show_col_types = FALSE)
test4_sub<-test4[1:10]
colnames(test4_sub)<-c("ID","X","Y","Map X","Map Y", "Lat","Lon","B1","B2","B3")
test5<-read_table("orthovnir069_R0I_Blue_Tarps.txt",show_col_types = FALSE)
test5_sub<-test5[1:10]
colnames(test5_sub)<-c("ID","X","Y","Map X","Map Y", "Lat","Lon","B1","B2","B3")
test6<-read_table("orthovnir069_R0I_NOT_Blue_Tarps.txt",show_col_types = FALSE)
test6_sub<-test6[1:10]
colnames(test6_sub)<-c("ID","X","Y","Map X","Map Y", "Lat","Lon","B1","B2","B3")
test7<-read_table("orthovnir078_R0I_Blue_Tarps.txt",show_col_types = FALSE)
test7_sub<-test7[1:10]
colnames(test7_sub)<-c("ID","X","Y","Map X","Map Y", "Lat","Lon","B1","B2","B3")
test8<-read_table("orthovnir078_R0I_NON_Blue_Tarps.txt",show_col_types = FALSE)
test8_sub<-test8[1:10]
colnames(test8_sub)<-c("ID","X","Y","Map X","Map Y", "Lat","Lon","B1","B2","B3")

combined_df_holdout<-bind_rows(test1_sub,test3_sub,test4_sub,test5_sub,test6_sub,test7_sub,
test8_sub)
df_fig<-combined_df_holdout[8:10]
df_ID<-combined_df_holdout[1]
df_figID<-cbind(df_ID,df_fig)
lonlatcords<-combined_df_holdout[6:10]
trainwoclass<-train[2:4]

```

Pre-processing

```

NonBlueholdout<-bind_rows(test1_sub,test4_sub,test6_sub,test7_sub,test8_sub)
NonBlueholdout_new<-NonBlueholdout[8:10]
names(NonBlueholdout_new)[1]<-"Red"
names(NonBlueholdout_new)[2]<-"Green"
names(NonBlueholdout_new)[3]<-"Blue"
NonBlueholdout_new <- NonBlueholdout_new %>%
  mutate(Class=1)
NonBlueholdout_new <- NonBlueholdout_new %>%
  mutate(Class=factor(Class,labels=c("1")))%>%
  mutate(Class=case_when(
    Class == 1 ~ "Non-Blue_Tarp"
  ))
Blueholdout<-bind_rows(test2,test3_sub,test5_sub)
Blueholdout_new<-Blueholdout[1:3]
names(Blueholdout_new)[1]<-"Red"
names(Blueholdout_new)[2]<-"Green"
names(Blueholdout_new)[3]<-"Blue"
Blueholdout_new <- Blueholdout_new %>%
  mutate(Class=1)
Blueholdout_new <- Blueholdout_new %>%
  mutate(Class=factor(Class,labels=c("1")))%>%
  mutate(Class=case_when(
    Class == 1 ~ "Blue_Tarp"
  ))
holdout<-bind_rows(Blueholdout_new,NonBlueholdout_new)

```

```

holdout<-holdout %>%
  mutate(type=case_when(
    Class == "Non-Blue_Tarp" ~ 0,
    Class == "Blue_Tarp" ~ 1
  ))

```

```

holdout<-holdout %>%
  mutate(
    group = paste(type, Class, sep="_"),
    group = factor(group),
  )

```

Training

```
Haiti_train<-train %>%
  mutate(type=case_when(
    Class == "Rooftop" ~ "Non-Blue_Tarp",
    Class == "Soil" ~ "Non-Blue_Tarp",
    Class == "Various Non-Tarp" ~ "Non-Blue_Tarp",
    Class == "Vegetation" ~ "Non-Blue_Tarp",
    Class == "Blue Tarp" ~ "Blue_Tarp"
  ))%>%
  mutate(Class=case_when(
    Class == "Rooftop" ~ "Non-Blue_Tarp",
    Class == "Soil" ~ "Non-Blue_Tarp",
    Class == "Various Non-Tarp" ~ "Non-Blue_Tarp",
    Class == "Vegetation" ~ "Non-Blue_Tarp",
    Class == "Blue Tarp" ~ "Blue_Tarp"
  ))%>%
  mutate(
    type=factor(type, levels=c("Blue_Tarp", "Non-Blue_Tarp")))%>%
  mutate(
    group = paste(type, Class, sep="_"),
    group = factor(group),
  )
set.seed(1)
formula<-Class~Red + Green + Blue
Haiti_recipe <- recipe(formula, data=Haiti_train) %>%
  step_normalize(all_numeric_predictors())
names(df_fig)[1]<-"Red"
names(df_fig)[2]<-"Green"
names(df_fig)[3]<-"Blue"
```

Logistic, LDA, and QDA Models

```
holdout<-holdout%>%
  mutate(
    Class=factor(Class))
Haiti_train<-Haiti_train%>%
  mutate(
    Class=factor(Class))
logreg_model_full <- logistic_reg(mode="classification", engine="glm") %>%
  fit(formula, Haiti_train)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
lda_model <- discrim_linear(mode="classification", engine="MASS") %>%
  fit(formula, Haiti_train)
qda_model <- discrim_quad(mode="classification", engine="MASS") %>%
  fit(formula,Haiti_train)
```

```
resamples <- vfold_cv(holdout, v=10, strata=type)
custom_metrics <- metric_set(roc_auc, accuracy)
cv_control <- control_resamples(save_pred=TRUE)
```

```
calculate_metrics <- function(model, train, test, model_name) {
  bind_rows(
    # Accuracy of training set
    bind_cols(
      model=model_name,
      dataset="train",
      metrics(model %>% augment(Haiti_train), truth=Class, estimate=.pred_class),
    ),
    # AUC of ROC curve of training set
    bind_cols(
      model=model_name,
      dataset="train",
      roc_auc(model %>% augment(Haiti_train), Class, .pred_Blue_Tarp, event_level="first"),
    ),
    # Accuracy of holdout set
    bind_cols(
      model=model_name,
      dataset="test",
      metrics(model %>% augment(holdout), truth=Class, estimate=.pred_class),
    ),
    # AUC of ROC curve of holdout set
    bind_cols(
      model=model_name,
      dataset="test",
      roc_auc(model %>% augment(holdout), Class, .pred_Blue_Tarp, event_level="first"),
    ),
  )
}
```

```
metrics_table <- function(all_metrics, caption) {
  all_metrics <- all_metrics %>% arrange(model, desc(dataset))
  all_metrics %>%
    pivot_wider(names_from=.metric, values_from=.estimate) %>%
    dplyr::select(-.estimator) %>%
    knitr::kable(caption=caption, digits=5) %>%
    kableExtra::kable_styling(full_width=FALSE)
}
```

```
all_metrics <- bind_rows(
  calculate_metrics(logreg_model_full, Haiti_train, holdout, "Logistic Regression"),
  calculate_metrics(lda_model, Haiti_train, holdout, "LDA"),
  calculate_metrics(qda_model, Haiti_train, holdout, "QDA"),
)
metrics_table(all_metrics, "Metrics for the classification models")
```

Metrics for the classification models

model	dataset	accuracy	kap	roc_auc
LDA	train	0.98397	0.75336	0.98888
LDA	test	0.98014	0.38800	0.99146
Logistic Regression	train	0.99529	0.92073	0.99851
Logistic Regression	test	0.98817	0.56108	0.99840
QDA	train	0.99461	0.90604	0.99822
QDA	test	0.99491	0.68267	0.99209

```


get_roc_plot <- function(model, data, model_name) {
  roc_data <- model %>%
  augment(Haiti_train) %>%
  roc_curve(truth=Class, .pred_Blue_Tarp, event_level="first")
  roc_auc(model %>% augment(Haiti_train), Class, .pred_Blue_Tarp, event_level="first")
  g <- autoplot(roc_data) +
  labs(title=model_name)
  return(g)
}

```

```


bind_rows(
  augment(logreg_model_full, Haiti_train) %>% mutate(model="Logistic Regression Training Set"),
  augment(lda_model, Haiti_train) %>% mutate(model="LDA Training Set"),
  augment(qda_model, Haiti_train) %>% mutate(model="QDA Training Set")
) %>%
group_by(model) %>%
roc_curve(truth=Class, .pred_Blue_Tarp, event_level="first") %>%
autoplot() + labs(title="ROC Plots Training Set")

```




```
get_roc_plot <- function(model, data, model_name) {
  roc_data <- model %>%
    augment(holdout) %>%
    roc_curve(truth=Class, .pred_Blue_Tarp, event_level="first")
  roc_auc(model %>% augment(holdout), Class, .pred_Blue_Tarp, event_level="first")
  g <- autoplot(roc_data) +
    labs(title=model_name)
  return(g)
}
```

```
bind_rows(  
  augment(logreg_model_full, holdout) %>% mutate(model="Logistic Regression Holdout Set"),  
  augment(lda_model, holdout) %>% mutate(model="LDA Holdout Set"),  
  augment(qda_model, holdout) %>% mutate(model="QDA Holdout Set")  
) %>%  
group_by(model) %>%  
roc_curve(truth=Class, .pred_Blue_Tarp, event_level="first") %>%  
autoplot() + labs(title="ROC Plots Holdout Set")
```





```
threshold_scan <- function(model, data, model_name) {  
  threshold_data <- model %>%  
  augment(holdout) %>%  
  probably::threshold_perf(Class, .pred_Blue_Tarp,  
    thresholds=seq(0.05, 0.95, 0.01), event_level="first",  
    metrics=metric_set(f_meas))  
  opt_threshold <- threshold_data %>%  
  arrange(-.estimate) %>%  
  first()  
  thresh <- ggplot(threshold_data, aes(x=.threshold, y=.estimate)) +  
    geom_line() +  
    geom_point(data=opt_threshold, color="red", size=2) +  
    labs(title=model_name, x="Threshold", y="F-Measure") +  
    coord_cartesian(ylim=c(0.3, 0.9))  
  return(list(  
    graph=thresh,  
    threshold=opt_threshold %>%  
    pull(.threshold)  
  ))  
}  
thresh1 <- threshold_scan(logreg_model_full, test, "Logistic regression")  
thresh2 <- threshold_scan(lda_model, test, "LDA")  
thresh3 <- threshold_scan(qda_model, test, "QDA")  
logreg_threshold <- thresh1$threshold  
lda_threshold <- thresh2$threshold  
qda_threshold <- thresh3$threshold  
thresh1$graph + thresh2$graph + thresh3$graph
```


Logistic regression



LDA



QDA



```
thresh1$threshold
```

```
## [1] 0.95
```

```
thresh2$threshold
```

```
## [1] 0.73
```

```
thresh3$threshold
```

```
## [1] 0.42
```

```
logthresh<-0.95
ldathresh<-0.73
qdathresh<-0.42
```

```

predict_at_threshold <- function(model, data, threshold) {
  return(
    model %>%
    augment(holdout) %>%
    mutate(.pred_class = make_two_class_pred(.pred_Blue_Tarp,
      c("Blue_Tarp", "Non-Blue_Tarp"), threshold=threshold)
    )
  )
}

predictions_logreg <- predict_at_threshold(logreg_model_full, holdout, logreg_threshold)
predictions_lda <- predict_at_threshold(lda_model, holdout, lda_threshold)
predictions_qda <- predict_at_threshold(qda_model, holdout, qda_threshold)

conf_mat(predictions_logreg, truth=Class, estimate=.pred_class)

```

```

##              Truth
## Prediction    Blue_Tarp Non-Blue_Tarp
##   Blue_Tarp      14991         7069
##   Non-Blue_Tarp    729        1985834

```

```

conf_mat(predictions_lda, truth=Class, estimate=.pred_class)

```

```

##              Truth
## Prediction    Blue_Tarp Non-Blue_Tarp
##   Blue_Tarp      12466         33883
##   Non-Blue_Tarp   3254        1959020

```

```

conf_mat(predictions_qda, truth=Class, estimate=.pred_class)

```

```

##              Truth
## Prediction    Blue_Tarp Non-Blue_Tarp
##   Blue_Tarp      11444         6164
##   Non-Blue_Tarp   4276        1986739

```

```

metrics<-metric_set(yardstick::accuracy,yardstick::sens,yardstick::spec,yardstick::f_measures,yardstick::j_index)

```

```

calculate_metrics_at_threshold <- function(model, test, model_name, threshold) {
  bind_rows(
    bind_cols(
      model=model_name, dataset="test", threshold=threshold,
      metrics(predict_at_threshold(model, holdout, threshold),
        truth=Class, estimate=.pred_class),
    ),
  )
}

metrics_at_threshold <- bind_rows(
  calculate_metrics_at_threshold(logreg_model_full, holdout, "Logistic regression", lo
gthresh),
  calculate_metrics_at_threshold(lda_model, holdout, "LDA", ldathresh),
  calculate_metrics_at_threshold(qda_model, holdout, "QDA", qdathresh),
) %>% arrange(dataset)

metrics_table(metrics_at_threshold, "Performance metrics with optimized threshold")

```

Performance metrics with optimized threshold

model	dataset	threshold	accuracy	sens	spec	f_meas	j_index
LDA	test	0.73	0.98151	0.79300	0.98300	0.40168	0.77600
Logistic regression	test	0.95	0.99612	0.95363	0.99645	0.79359	0.95008
QDA	test	0.42	0.99480	0.72799	0.99691	0.68675	0.72490

```
Haiti_train<-train %>%
  mutate(type=case_when(
    Class == "Rooftop" ~ "Non-Blue_Tarp",
    Class == "Soil" ~ "Non-Blue_Tarp",
    Class == "Various Non-Tarp" ~ "Non-Blue_Tarp",
    Class == "Vegetation" ~ "Non-Blue_Tarp",
    Class == "Blue Tarp" ~ "Blue_Tarp"
  ))%>%
  mutate(Class=case_when(
    Class == "Rooftop" ~ "Non-Blue_Tarp",
    Class == "Soil" ~ "Non-Blue_Tarp",
    Class == "Various Non-Tarp" ~ "Non-Blue_Tarp",
    Class == "Vegetation" ~ "Non-Blue_Tarp",
    Class == "Blue Tarp" ~ "Blue_Tarp"
  ))%>%
  mutate(
    type=factor(type, levels=c("Blue_Tarp", "Non-Blue_Tarp"))%>%
    mutate(
      group = paste(type, Class, sep="_"),
      group = factor(group),
    )
  )
set.seed(1)
formula<-Class~Red + Green + Blue
Haiti_recipe <- recipe(formula, data=Haiti_train) %>%
  step_normalize(all_numeric_predictors())
tuningtrain<-Haiti_train[1:4]
tuningtrain_sub<- tuningtrain %>%
  mutate(Class=factor(Class,labels=c("Non-Blue_Tarp","Blue_Tarp")))%>%
  mutate(case_when(
    Class == "Blue_Tarp" ~ 1,
    Class == "Non-Blue_Tarp" ~ 0
  ))
```

Random Forests:

```
tuningtrain$Red<-as.numeric(tuningtrain$Red)
tuningtrain$Green<-as.numeric(tuningtrain$Green)
tuningtrain$Blue<-as.numeric(tuningtrain$Blue)
tuningtrain$Class<-as.character(tuningtrain$Class)
formula<-Class~ Red + Green + Blue
rec <- recipe(formula, data=tuningtrain)
class(tuningtrain$Class)
```

```
## [1] "character"
```


```
forest<-rand_forest(mode="classification", trees=500,min_n = tune(), mtry=tune()) %>%
  set_engine("ranger",importance="impurity")

forests_workflow <- workflow() %>%
  add_recipe(rec) %>%
  add_model(forest)
```

```
parameters <- extract_parameter_set_dials(forests_workflow)%>%
  update(
    mtry = mtry(c(1, 3)),
    min_n = min_n(c(2, 40))
  )
```

```
set.seed(1)
tune_results_forest <- tune_grid(forests_workflow,
  resamples=vfold_cv(tuningtrain),
  grid=grid_regular(parameters))
```

```
autoplot(tune_results_forest)
```



```
show_best(tune_results_forest,metric="roc_auc")
```

```
## # A tibble: 5 × 8
##   mtry min_n .metric .estimator mean      n std_err .config
##   <int> <int> <chr>    <chr>    <dbl> <int>    <dbl> <chr>
## 1     2    21 roc_auc binary     1.00    10 0.0000348 Preprocessor1_Model5
## 2     2     2 roc_auc binary     1.00    10 0.0000344 Preprocessor1_Model2
## 3     2    40 roc_auc binary     1.00    10 0.0000388 Preprocessor1_Model8
## 4     1    21 roc_auc binary     1.00    10 0.0000401 Preprocessor1_Model4
## 5     1    40 roc_auc binary     1.00    10 0.0000383 Preprocessor1_Model7
```

```
best_parameters <- select_best(tune_results_forest, metric="roc_auc")
best_workflow <- forests_workflow %>%
  finalize_workflow(best_parameters) %>%
  fit(tuningtrain)
best_workflow
```

```
## == Workflow [trained] ==
## Preprocessor: Recipe
## Model: rand_forest()
##
## == Preprocessor ==
## 0 Recipe Steps
##
## == Model ==
## Ranger result
##
## Call:
## ranger::ranger(x = maybe_data_frame(x), y = y, mtry = min_cols(~2L, x), num.trees = ~500, min.node.size = min_rows(~21L, x), importance = ~"impurity", num.threads = 1, verbose = FALSE, seed = sample.int(10^5, 1), probability = TRUE)
##
## Type: Probability estimation
## Number of trees: 500
## Sample size: 63241
## Number of independent variables: 3
## Mtry: 2
## Target node size: 21
## Variable importance mode: impurity
## Splitrule: gini
## OOB prediction error (Brier s.): 0.002167524
```

```
resamples <- vfold_cv(tuningtrain, v=10, strata=Class)
custom_metrics <- metric_set(roc_auc, accuracy)
cv_control <- control_resamples(save_pred=TRUE)


forest_cv <- fit_resamples(best_workflow, resamples, metrics=custom_metrics, control=cv_control)
```

```
cv_metrics <- bind_rows(  
  collect_metrics(forest_cv) %>%  
    mutate(model="Random Forest"),  
)  
cv_metrics %>%  
  dplyr::select(model, .metric, mean) %>%  
  pivot_wider(names_from=.metric, values_from=mean) %>%  
  knitr::kable(caption="Random Forest Metrics", digits=5)
```

Random Forest Metrics

model	accuracy	roc_auc
Random Forest	0.99703	0.99974

```
bind_rows(  
  collect_predictions(forest_cv) %>% mutate(model="Random Forest"),  
) %>%  
  group_by(model) %>%  
  roc_curve(truth=Class, .pred_Blue_Tarp, event_level="first") %>%  
  autoplot()+labs(title="Random Forest ROC plot")
```

```
holdout<-holdout%>%
  mutate(
    Class=as.character(Class))
result_cvh <- fit_resamples(best_workflow, resamples,
                             metrics=custom_metrics, control=cv_control)
fitted_model <- best_workflow %>% fit(holdout)
```

```

cv_metrics <- bind_rows(
  collect_metrics(result_cvh) %>%
    mutate(model="Random Forests"),
)
cv_metrics %>%
  dplyr::select(model, .metric, mean) %>%
  pivot_wider(names_from=.metric, values_from=mean) %>%
  knitr::kable(caption="Random Forests Metrics", digits=5)

```

Random Forests Metrics


model	accuracy	roc_auc
Random Forests	0.99693	0.99974

```

holdout<-holdout%>%
  mutate(
    Class=factor(Class))
threshold_scan <- function(model, data, model_name) {
  threshold_data <- model %>%
  augment(holdout) %>%
  probably::threshold_perf(Class, .pred_Blue_Tarp,
  thresholds=seq(0.05, 0.95, 0.01), event_level="first",
  metrics=metric_set(f_meas))
  opt_threshold <- threshold_data %>%
  arrange(-.estimate) %>%
  first()
  thresh <- ggplot(threshold_data, aes(x=.threshold, y=.estimate)) +
  geom_line() +
  geom_point(data=opt_threshold, color="red", size=2) +
  labs(title=model_name, x="Threshold", y="F-Measure") +
  coord_cartesian(ylim=c(0.3, 1.0))
  return(list(
    graph=thresh,
    threshold=opt_threshold %>%
    pull(.threshold)
  ))
}
thresh4 <- threshold_scan(fitted_model, test, "Random Forest")
Forest_threshold <- thresh4$threshold
thresh4$graph

```

Random Forest



```
thresh4$threshold
```

```
## [1] 0.6
```

```
predict_at_threshold <- function(model, data, threshold) {
  return(
    model %>%
    augment(holdout) %>%
    mutate(.pred_class = make_two_class_pred(.pred_Blue_Tarp,
      c("Blue_Tarp", "Non-Blue_Tarp"), threshold=threshold)
    )
  )
}

predictions_Forest <- predict_at_threshold(fitted_model, holdout, Forest_threshold)
conf_mat(predictions_Forest, truth=Class, estimate=.pred_class)
```

```
##           Truth
## Prediction  Blue_Tarp Non-Blue_Tarp
##   Blue_Tarp    14954      636
##   Non-Blue_Tarp    766    1992267
```

```
metrics<-metric_set(yardstick::accuracy,yardstick::sens,yardstick::spec,yardstick::f_meas,
yardstick::j_index)
```

```
calculate_metrics_at_threshold <- function(model, test, model_name, threshold) {
  bind_rows(
    bind_cols(
      model=model_name, dataset="test", threshold=threshold,
      metrics(predict_at_threshold(model, holdout, threshold),
        truth=Class, estimate=.pred_class),
    ),
  )
}

metrics_at_threshold <- bind_rows(
  calculate_metrics_at_threshold(logreg_model_full,holdout, "Logistic regression", log
thresh),
  calculate_metrics_at_threshold(lda_model, holdout, "LDA", ldathresh),
  calculate_metrics_at_threshold(qda_model, holdout, "QDA", qdathresh),
  calculate_metrics_at_threshold(fitted_model, holdout, "Random Forest", Forest_thresh
old),
) %>% arrange(dataset)

metrics_table(metrics_at_threshold, "Performance metrics with optimized threshold")
```

Performance metrics with optimized threshold

model	dataset	threshold	accuracy	sens	spec	f_meas	j_index
LDA	test	0.73	0.98151	0.79300	0.98300	0.40168	0.77600
Logistic regression	test	0.95	0.99612	0.95363	0.99645	0.79359	0.95008
QDA	test	0.42	0.99480	0.72799	0.99691	0.68675	0.72490
Random Forest	test	0.60	0.99930	0.95127	0.99968	0.95522	0.95095


Xg Boost

```
Haiti_train<-Haiti_train%>%
  mutate(
    Class=as.character(Class))
library(bonsai) # this is required
boost_wf <- workflow() %>%
add_recipe(recipe(formula, data=Haiti_train)) %>%
add_model(boost_tree(mode="classification", engine="lightgbm",
trees=500, tree_depth=tune(), learn_rate=tune()))
```

```
parameters <- extract_parameter_set_dials(boost_wf)
tune_xgboost <- tune_bayes(boost_wf,
  resamples=resamples,
  metrics=custom_metrics,
  param_info=parameters, iter=25)
```

```
## ! No improvement for 10 iterations; returning current results.
```

```
autoplot(tune_xgboost)
```



```
select_best(tune_xgboost, metric='roc_auc')
```

```
## # A tibble: 1 × 3
##   tree_depth learn_rate .config
##       <int>      <dbl> <chr>
## 1         5    0.0734 Iter8
```

```
holdout<-holdout%>%
  mutate(
    Class=as.character(Class))
best_boost_wf <- boost_wf %>%
  finalize_workflow(select_best(tune_xgboost, metric='roc_auc'))
boost_model<-best_boost_wf %>% fit(holdout)
```

```
boost_cv <- fit_resamples(best_boost_wf, resamples, metrics=custom_metrics, control=cv_c
ontrol)
```


```
cv_metrics <- bind_rows(
  collect_metrics(boost_cv) %>%
    mutate(model="Xg Boost"),
)
cv_metrics %>%
  dplyr::select(model, .metric, mean) %>%
  pivot_wider(names_from=.metric, values_from=mean) %>%
  knitr::kable(caption="Xg Boost Metrics", digits=5)
```

Xg Boost Metrics


model	accuracy	roc_auc
Xg Boost	0.99685	0.99945

```
roc_cv_plot <- function(model_cv, model_name) {
  cv_predictions <- collect_predictions(model_cv)
  cv_roc <- cv_predictions %>%
    roc_curve(truth=Class, .pred_Blue_Tarp, event_level="first")
return(autoplot(cv_roc) + labs(title=model_name))
}
```

```
roc_boost <- roc_cv_plot(boost_cv, "Boosting")
roc_boost
```



```
holdout<-holdout%>%
  mutate(
    Class=factor(Class))
threshold_scan <- function(model, data, model_name) {
  threshold_data <- model %>%
  augment(holdout) %>%
  probably::threshold_perf(Class, .pred_Blue_Tarp,
  thresholds=seq(0.05, 0.95, 0.01), event_level="first",
  metrics=metric_set(f_meas))
  opt_threshold <- threshold_data %>%
  arrange(-.estimate) %>%
  first()
  thresh <- ggplot(threshold_data, aes(x=.threshold, y=.estimate)) +
  geom_line() +
  geom_point(data=opt_threshold, color="red", size=2) +
  labs(title=model_name, x="Threshold", y="F-Measure") +
  coord_cartesian(ylim=c(0.3, 1.0))
  return(list(
  graph=thresh,
  threshold=opt_threshold %>%
  pull(.threshold)
  ))
}
thresh5 <- threshold_scan(boost_model, test, "Xg Boost")
Boost_threshold <- thresh5$threshold
thresh5$graph
```

```
thresh5$threshold
```

```
## [1] 0.46
```

```
predict_at_threshold <- function(model, data, threshold) {
  return(
    model %>%
    augment(holdout) %>%
    mutate(.pred_class = make_two_class_pred(.pred_Blue_Tarp,
      c("Blue_Tarp", "Non-Blue_Tarp"), threshold=threshold)
    )
  )
}

predictions_Boost <- predict_at_threshold(boost_model, holdout, Boost_threshold)
conf_mat(predictions_Boost, truth=Class, estimate=.pred_class)
```

```
##           Truth
## Prediction  Blue_Tarp Non-Blue_Tarp
##   Blue_Tarp    15258      1370
##   Non-Blue_Tarp   462      1991533
```

```
metrics<-metric_set(yardstick::accuracy,yardstick::sens,yardstick::spec,yardstick::f_meas,
yardstick::j_index)
```

```
calculate_metrics_at_threshold <- function(model, test, model_name, threshold) {
  bind_rows(
    bind_cols(
      model=model_name, dataset="test", threshold=threshold,
      metrics(predict_at_threshold(model, holdout, threshold),
        truth=Class, estimate=.pred_class),
    ),
  )
}

metrics_at_threshold <- bind_rows(
  calculate_metrics_at_threshold(logreg_model_full,holdout, "Logistic regression", log
thresh),
  calculate_metrics_at_threshold(lda_model, holdout, "LDA", ldathresh),
  calculate_metrics_at_threshold(qda_model, holdout, "QDA", qdathresh),
  calculate_metrics_at_threshold(fitted_model, holdout, "Random Forest", Forest_thresh
old),
  calculate_metrics_at_threshold(boost_model, holdout, "Xg Boost", Boost_threshold),
) %>% arrange(dataset)

metrics_table(metrics_at_threshold, "Performance metrics with optimized threshold")
```

Performance metrics with optimized threshold

model	dataset	threshold	accuracy	sens	spec	f_meas	j_index
LDA	test	0.73	0.98151	0.79300	0.98300	0.40168	0.77600
Logistic regression	test	0.95	0.99612	0.95363	0.99645	0.79359	0.95008
QDA	test	0.42	0.99480	0.72799	0.99691	0.68675	0.72490
Random Forest	test	0.60	0.99930	0.95127	0.99968	0.95522	0.95095
Xg Boost	test	0.46	0.99909	0.97061	0.99931	0.94337	0.96992

Penalized Logistic Regression (elastic net penalty):

```
Haiti_train<-Haiti_train%>%
  mutate(
    Class=as.character(Class))
formula<-Class~Red+Green+Blue

recipe_spec <- recipe(formula, data=Haiti_train) %>%
  step_dummy(all_nominal(), -all_outcomes())
```


```
model_glmnet <- logistic_reg(engine="glmnet", mode="classification",
                             penalty=tune(), mixture=tune())
wf_penalized <- workflow() %>%
  add_model(model_glmnet) %>%
  add_recipe(recipe_spec)
```

```
parameters <- extract_parameter_set_dials(wf_penalized) %>%
  update(penalty=penalty(c(-4, -1)))
tune_results_penalized <- tune_grid(wf_penalized,
                                    resamples=vfold_cv(Haiti_train),
                                    grid=grid_regular(parameters))
```

```
tune_wf_penalized <- tune_bayes(wf_penalized, resamples=resamples, metrics=custom_metrics,
                                param_info=parameters, iter=25)
```

```
## ! No improvement for 10 iterations; returning current results.
```

```
autoplot(tune_wf_penalized)
```



```
select_best(tune_wf_penalized, metric="roc_auc")
```

```
## # A tibble: 1 × 3
##   penalty mixture .config
##   <dbl>    <dbl> <chr>
## 1 0.000119  0.901 Iter7
```

```
best_parameter_penalized <- select_best(tune_wf_penalized, metric="roc_auc")
best_wf_penalized <- finalize_workflow(wf_penalized, best_parameter_penalized)
```

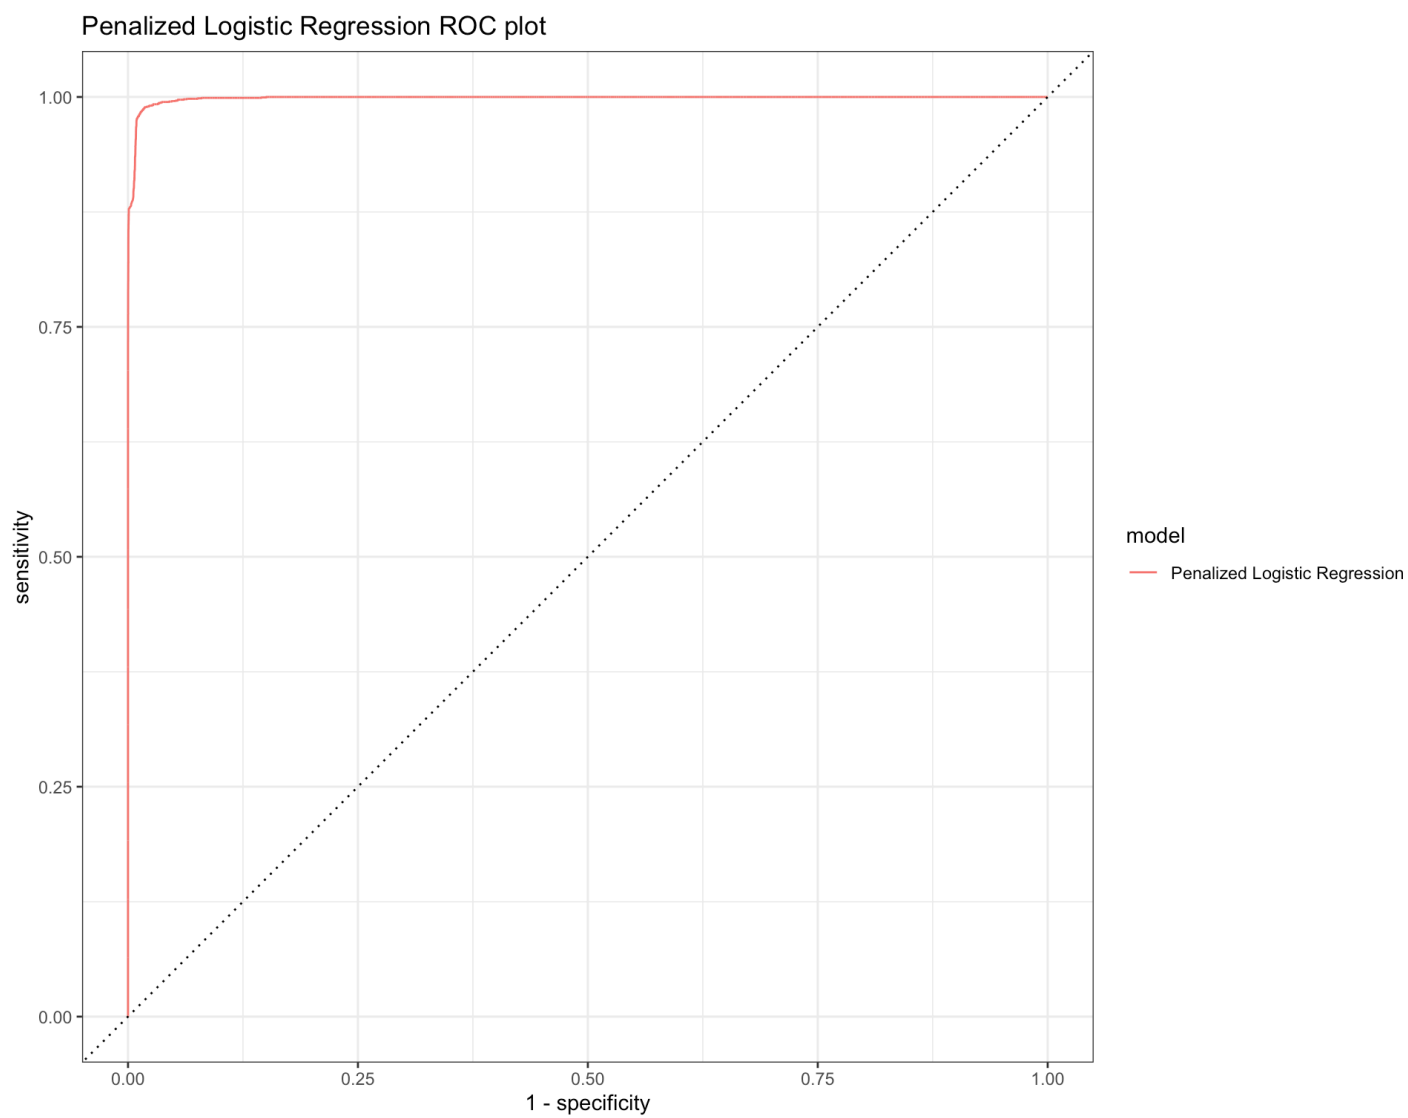
```
holdout<-holdout%>%
  mutate(
    Class=as.character(Class))
result_cv_penalized <- fit_resamples(best_wf_penalized, resamples,
                                     metrics=custom_metrics, control=cv_control)
fitted_model_penalized <- best_wf_penalized %>% fit(holdout)
```

```
cv_metrics <- bind_rows(
  collect_metrics(result_cv_penalized) %>%
    mutate(model="Penalized Logistic Regression"),
)
cv_metrics %>%
  dplyr::select(model, .metric, mean) %>%
  pivot_wider(names_from=.metric, values_from=mean) %>%
  knitr::kable(caption="Penalized Logistic Regression Metrics", digits=5)
```

Penalized Logistic Regression Metrics


model	accuracy	roc_auc
Penalized Logistic Regression	0.9948	0.9985

```
bind_rows(
  collect_predictions(result_cv_penalized) %>% mutate(model="Penalized Logistic Regression"),
) %>%
  group_by(model) %>%
  roc_curve(truth=Class, .pred_Blue_Tarp, event_level="first") %>%
  autoplot()+labs(title="Penalized Logistic Regression ROC plot")
```



```
holdout<-holdout%>%
  mutate(
    Class=factor(Class))
threshold_scan <- function(model, data, model_name) {
  threshold_data <- model %>%
  augment(holdout) %>%
  probably::threshold_perf(Class, .pred_Blue_Tarp,
  thresholds=seq(0.05, 0.95, 0.01), event_level="first",
  metrics=metric_set(f_meas))
  opt_threshold <- threshold_data %>%
  arrange(-.estimate) %>%
  first()
  thresh <- ggplot(threshold_data, aes(x=.threshold, y=.estimate)) +
  geom_line() +
  geom_point(data=opt_threshold, color="red", size=2) +
  labs(title=model_name, x="Threshold", y="F-Measure") +
  coord_cartesian(ylim=c(0.3, 1.0))
  return(list(
  graph=thresh,
  threshold=opt_threshold %>%
  pull(.threshold)
  ))
}
thresh6 <- threshold_scan(fitted_model_penalized, test, "Penalized Logistic Regression")
Penalized_threshold <- thresh6$threshold
thresh6$graph
```

Penalized Logistic Regression



```
thresh6$threshold
```

```
## [1] 0.08
```

```
predict_at_threshold <- function(model, data, threshold) {
  return(
    model %>%
    augment(holdout) %>%
    mutate(.pred_class = make_two_class_pred(.pred_Blue_Tarp,
      c("Blue_Tarp", "Non-Blue_Tarp"), threshold=threshold)
    )
  )
}

predictions_Penalized <- predict_at_threshold(fitted_model_penalized, holdout, Penalized_
_threshold)
conf_mat(predictions_Penalized, truth=Class, estimate=.pred_class)
```

```
##           Truth
## Prediction  Blue_Tarp Non-Blue_Tarp
##   Blue_Tarp   13851      4541
##   Non-Blue_Tarp  1869     1988362
```

```
metrics<-metric_set(yardstick::accuracy,yardstick::sens,yardstick::spec,yardstick::f_measures,yardstick::j_index)
```

```
calculate_metrics_at_threshold <- function(model, test, model_name, threshold) {
  bind_rows(
    bind_cols(
      model=model_name, dataset="test", threshold=threshold,
      metrics(predict_at_threshold(model, holdout, threshold),
        truth=Class, estimate=.pred_class),
    ),
  )
}

metrics_at_threshold <- bind_rows(
  calculate_metrics_at_threshold(logreg_model_full,holdout, "Logistic regression", logthresh),
  calculate_metrics_at_threshold(lda_model, holdout, "LDA", ldathresh),
  calculate_metrics_at_threshold(qda_model, holdout, "QDA", qdathresh),
  calculate_metrics_at_threshold(fitted_model, holdout, "Random Forest", Forest_threshold),
  calculate_metrics_at_threshold(boost_model, holdout, "Xg Boost", Boost_threshold),
  calculate_metrics_at_threshold(fitted_model_penalized, holdout, "Penalized Logistic Regression", Penalized_threshold),
) %>% arrange(dataset)

metrics_table(metrics_at_threshold, "Performance metrics with optimized threshold")
```

Performance metrics with optimized threshold

model	dataset	threshold	accuracy	sens	spec	f_measures	j_index
LDA	test	0.73	0.98151	0.79300	0.98300	0.40168	0.77600
Logistic regression	test	0.95	0.99612	0.95363	0.99645	0.79359	0.95008
Penalized Logistic Regression	test	0.08	0.99681	0.88111	0.99772	0.81209	0.87883
QDA	test	0.42	0.99480	0.72799	0.99691	0.68675	0.72490
Random Forest	test	0.60	0.99930	0.95127	0.99968	0.95522	0.95095
Xg Boost	test	0.46	0.99909	0.97061	0.99931	0.94337	0.96992

All Confusion matrices

```
conf_mat(predictions_logreg, truth=Class, estimate=.pred_class)
```



```
##                Truth
## Prediction      Blue_Tarp Non-Blue_Tarp
##   Blue_Tarp      14991      7069
##   Non-Blue_Tarp    729      1985834
```

```
conf_mat(predictions_lda, truth=Class, estimate=.pred_class)
```

```
##                Truth
## Prediction      Blue_Tarp Non-Blue_Tarp
##   Blue_Tarp      12466      33883
##   Non-Blue_Tarp    3254      1959020
```

```
conf_mat(predictions_qda, truth=Class, estimate=.pred_class)
```

```
##                Truth
## Prediction      Blue_Tarp Non-Blue_Tarp
##   Blue_Tarp      11444      6164
##   Non-Blue_Tarp    4276      1986739
```

```
conf_mat(predictions_Forest, truth=Class, estimate=.pred_class)
```

```
##                Truth
## Prediction      Blue_Tarp Non-Blue_Tarp
##   Blue_Tarp      14954      636
##   Non-Blue_Tarp    766      1992267
```

```
conf_mat(predictions_Boost, truth=Class, estimate=.pred_class)
```

```
##                Truth
## Prediction      Blue_Tarp Non-Blue_Tarp
##   Blue_Tarp      15258      1370
##   Non-Blue_Tarp    462      1991533
```

```
conf_mat(predictions_Penalized, truth=Class, estimate=.pred_class)
```

```
##                Truth
## Prediction      Blue_Tarp Non-Blue_Tarp
##   Blue_Tarp      13851      4541
##   Non-Blue_Tarp    1869      1988362
```

All F-Measures combined

```

holdout<-holdout%>%
  mutate(
    Class=factor(Class))
threshold_scan <- function(model, data, model_name) {
threshold_data <- model %>%
augment(holdout) %>%
probably::threshold_perf(Class, .pred_Blue_Tarp,
thresholds=seq(0.05, 0.95, 0.01), event_level="first",
metrics=metric_set(f_meas))
opt_threshold <- threshold_data %>%
arrange(-.estimate) %>%
first()
  thresh <- ggplot(threshold_data, aes(x=.threshold, y=.estimate)) +
    geom_line() +
    geom_point(data=opt_threshold, color="red", size=2) +
    labs(title=model_name, x="Threshold", y="F-Measure") +
    coord_cartesian(ylim=c(0.3, 1.0))
  return(list(
    graph=thresh,
    threshold=opt_threshold %>%
    pull(.threshold)
  ))
}
thresh1 <- threshold_scan(logreg_model_full, test, "Logistic regression")
thresh2 <- threshold_scan(lda_model, test, "LDA")
thresh3 <- threshold_scan(qda_model, test, "QDA")
thresh4 <- threshold_scan(fitted_model, test, "Random Forest")
thresh5 <- threshold_scan(boost_model, test, "Xg Boost")
thresh6 <- threshold_scan(fitted_model_penalized, test, "Penalized")

logreg_threshold <- thresh1$threshold
lda_threshold <- thresh2$threshold
qda_threshold <- thresh3$threshold
Forest_threshold <- thresh4$threshold
Boost_threshold <- thresh5$threshold

Penalized_threshold <- thresh6$threshold
thresh1$graph + thresh2$graph + thresh3$graph + thresh4$graph + thresh5$graph+ thresh6$g
raph

```

